

# Ruby to Wasm Compiler

Mark Murphy

Advisor: Brian Ladd

Semester: Fall 2024

September 8, 2025

## 1 Introduction

I propose to create an Ahead-of-Time (AoT) compiler for a subset of the Ruby programming language. The only compilation target will be WebAssembly 3.0, the newest version of WebAssembly, with support for garbage-collected memory [1].

WebAssembly (Wasm) is a virtual machine (VM), similar to the Java Virtual Machine or the .NET Common Language Runtime. Every major web browser includes a Wasm VM, allowing it to be used as a non-Javascript compilation target for compilers targeting the browser.

While WebAssembly has existed since 2017, the initial 1.0 release supported only manual memory management. Shipping Wasm 1.0 code for any garbage-collected language requires shipping a garbage collector. In contrast, Wasm 3.0 includes the WasmGC proposal [2], which uses the browser's existing garbage collector to implement managed data types. The WasmGC proposal is currently supported in all major browsers [3].

These changes greatly simplify the compilation of garbage-collected languages to Wasm, making it closer to a bytecode interpreter than a full compiler with native backend. There are WasmGC ports for Java, Kotlin, and OCaml. Of particular interest is Hoot, a compiler for Guile Scheme targeting WasmGC [4] [5]. Similarly to Ruby, Scheme is a dynamic, metaprogramming-heavy language, so this should be a good source.

There exist several projects which compile a Ruby interpreter to Wasm [6] [7]. However, to my knowledge there exists no project that attempts to compile Ruby code directly to WasmGC. This novelty is one reason I'm interested in Ruby in particular. I also understand Ruby to be somewhat more popular in Japan (its country of origin), which is a good fit for my future ambitions.

I intend to implement Ruby's object-oriented facilities, as well as the normal data structures expected of a high-level programming language. This will include implementations of strings (likely through Javascript interop), interned symbols, dynamically-sized lists, hash tables, and function closures. Because system and browser APIs aren't exposed directly to Wasm programs, implementing certain standard library features, including file handling and IO, would require Javascript library support.

I intend to write the compiler in Rust, because it's my favorite programming language. It's also a pretty convenient language for writing interpreters/compilers, since it has both ML-like algebraic data types and nice byte-manipulation facilities.

## 2 Deliverables

The main deliverable of this project is a compiler, which takes in files of the Ruby subset, and produces Wasm binaries. The compiler should also be able to .wat files, Wasm's textual format.

The project will also include a small demonstration program, which runs in the browser. Currently, I intend to produce a simple "to-do" application, which will demonstrate Javascript interop and DOM manipulation. That will consist of a Ruby file compiled to Wasm, a small HTML file, and a Javascript scaffolding file, which loads and runs the Wasm file, and provides the required library functions for browser interoperation.

I consider the project’s documentation to be a deliverable, as well as the capstone presentation slides. In addition to normal internal API documentation and instructions for installing and running the compiler, I will include fairly in-depth architectural documentation, explaining the compilation pipeline and discussing technical choices made.

### 3 Preparation

While this is certainly an ambitious project, I think several of my experiences in and out of class have prepared me for it.

Relevant classes include Theory of Computation, Operating Systems, Assembly Language, Programming Languages, and the always-relevant Professional Practice. While I didn’t submit completed versions of many steps of the Programming Languages scheme interpreter assignment, I did follow along with the project, and I feel I understand the techniques used.

I had previously intended to create a “fantasy console” for my capstone project, which would include a bytecode interpreter for a self-designed programming language. I eventually decided the scope of the project was to large, but in my attempts at prototyping, I completed a (buggy, unpleasant, incomplete) bytecode interpreter for a basic lisp.

I’ve made an emulator for the historical Chip 8 virtual console, which involved decoding a binary bytecode format.

I have submitted some minor contributions to Rust Analyzer, the Rust language server. I haven’t gone deeper than the textual and AST layers, but that experience will still be useful.

I have studied the Wasm 3.0 specification and secondary sources describing the purpose and technical specifics of the WasmGC proposal. I believe I have the understanding of the format required to complete this project.

### 4 Practice

I expect that this project will cover many of the topics traditionally covered in a Compilers class. Like a Compilers class, it will therefore tie together the theoretical, algorithmic, and concrete engineering aspects of computer science, and touch on material from nearly every other CS class (with the merciful exception of Networking). I will have the opportunity to review and expand upon the topics covered in those classes.

It will not interface with the operating system in the way that a normal compilers project would, but it will involve many of the low-level considerations brought up in the Operating Systems class.

Correct and understandable documentation of the project will require communication skills improved by several CS classes, particularly Software Engineering and Professional Practice. Preparation for the capstone presentation will similarly hone the in-person presentation skills we’ve covered.

### 5 Administration

The project will be completed in the Fall semester. I will register for 1 class, 3 credits.

I will break the project into one-week sprints. At the end of each sprint, I will produce a sprint report describing the work completed that sprint.

### A Grading rubric

I think the main axes of evaluation should be the quality of the compiler, the correctness of the compiler with respect to both Ruby and Wasm, the comprehensiveness of the implemented Ruby subset, quality of the documentation, quality of the demo, and quality of the presentation.

## **Rubric: 2.0 grade**

### **Compiler quality**

- Compiles a very small subset of Ruby programs, including basic statements, expressions and method declarations. Certain control structures, most data structures, and aspects of the object/class system are missing.
- Compiler successfully outputs Wasm programs for some correctly-formed programs
- Compiler crashes on many correctly-formed programs
- Compiler crashes without explanation on incorrect programs

### **Compiler correctness**

- Many compiled Wasm binaries fail to run in targeted browsers
- Semantics of compiled programs differ significantly from expected Ruby semantics
- Compiler may only emit .wat files, requiring external processing to convert to Wasm binary format

### **Documentation quality**

- Internal documentation is very incomplete, or is not inline with CS department coding guidelines
- Final compiler executable is undocumented or lightly documented. A user of the project would have to consult the project code to successfully run the project.
- No architectural explanation or justification of technical choices is included
- The available subset of Ruby is not documented or badly documented

### **Demo quality**

- Produces a demo program that demonstrates successful execution, but may do nothing useful, or produce only a return value rather than any visible side effects
- Demo runs in some Wasm runtime, but may only run on desktop or with special setup
- Demo program is clearly buggy, frequently crashing or demonstrating logic errors

### **Presentation quality**

- Presentation introduces the project, but doesn't describe it clearly, or leaves out important details
- Presenter demonstrates the functioning of the project in some way, but has significant technical challenges
- Presentation appears under-prepared. Presenter frequently pauses to consult notes
- Presenter is not able to respond effectively to audience questions

## Rubric: 3.0 grade

### Compiler quality

- Compiles a subset of Ruby programs, including most statements, expressions and method declarations. Most control structures, most data structures, and most of the object/class system are present.
- Compiler successfully outputs Wasm programs for most correctly-formed programs
- Compiler crashes on some correctly-formed programs
- Compiler gracefully terminates on incorrectly-formed programs, but does not produce many useful error messages

### Compiler correctness

- Some compiled Wasm binaries fail to run in targeted browsers
- Semantics of compiled programs differ minorly or infrequently from expected Ruby semantics
- Compiler emits Wasm bytecode

### Documentation quality

- Internal documentation is slightly incomplete, or differs somewhat CS department coding guidelines
- Final compiler executable is lightly documented. A user of the project might have to consult the project code for clarification on some
- Architectural explanation and justification of technical choices is included, but somewhat lacking
- The available subset of Ruby is documented, but may be unclear or missing some information

### Demo quality

- Demo runs in browser and produces observable side-effects
- Interoperation with Javascript is lacking
- Demo program is mostly correct, rarely crashing or demonstrating logic errors

### Presentation quality

- Presentation introduces the project, and describes it, but may assume too much prerequisite knowledge, or fail to be sufficiently detailed
- Presenter demonstrates the produced demo program, but may have minor technical problems
- Presentation appears reasonably well-prepared. Presenter may sometimes get stuck or consult notes
- Presenter is able to respond to many audience questions, but fails to answer some technical questions in a satisfying way

## Rubric: 4.0 grade

### Compiler quality

- Compiles a subset of Ruby that includes most or all control structures, many features of Ruby's object/class system, and many Ruby data structures
- Compiles almost all correctly-formed input programs
- Gracefully handles incorrect input programs
- Tracks source locations and provides basic error reporting in the lexer and parser

### **Compiler correctness**

- Ruby programs using only the documented subset of the language compile and run identically in mainstream Ruby interpreters and my compiler, with the exception of documented semantic differences and minor edge cases
- Produced programs compile and run without error in browser Wasm implementations

### **Documentation quality**

- Internal API documentation is complete and in line with CS department coding guidelines, perhaps with certain agreed-upon modifications
- Final compiler executable is documented from a user perspective, both in the project README and when running with the `-help` flag
- Documentation includes an architectural explanation, suitable for new project contributors with basic background knowledge. Architecture documentation includes explanation and justification of technical choices.
- The available subset of Ruby and semantic differences from mainstream Ruby interpreters are documented

### **Demo quality**

- Demo runs correctly in targeted browsers, demonstrating intended functionality without logic errors or browser-reported exceptions
- Demo makes non-trivial use of Ruby data structures
- Demo demonstrates Wasm-Javascript interoperation

### **Presentation quality**

- Presentation introduces project, rationale, and a technical/architectural overview
- Presentation demonstrates the operation of both the compiler executable and the demo project
- Presentation is delivered smoothly, and appears practiced
- Presenter responses to audience questions demonstrate understanding of the project and the underlying technology

## **B Tentative schedule**

I don't think the date of the Board of Advisors meeting has been announced. I'm assuming November 15, which was the date of the Fall 2024 BoA meeting.

I have open availability for meetings, and will update this document when a meeting schedule has been decided.

I intend to produce a brief weekly report, including a list of backlog items completed that week.

### **Week 1: 8/31 - 9/6**

- Install, compile and run reference programs – Hoot scheme compiler and a Ruby interpreter Play with program output and begin to understand the project structure, entry points, etc.
- Collect auxiliary tools – Wasm bintools, desktop Wasm interpreter, etc.
- Create scaffolding project to run and debug compiled Wasm files in Chrome

- Walking skeleton – end-to-end lexing, parsing, code generation for a few constant values and operators
- Start to pin down value representation. Look to the reference programs for help.

## Week 2: 9/7 - 9/13

- Finish the lexer or get close. Add basic line tracking and error reporting.
- Settle on an AST structure.
- Build out the parser.
- Settle on value representation, including immediate-value and heap integers.
- Implement control structures – if / else, while.

## Week 3: 9/14 - 9/20

- Local variable assignment and scoping. Pin down the Ruby scoping semantics.
- Global variable assignment. Wasm has facilities for both of these.
- Implement top-level method definitions.
- Remind myself how closure implementations work.
- Get a very clear understanding of Ruby's object system. See how the Ruby interpreters implement it.
- Research String implementation in Wasm.

## Week 4: 9/21 - 9/27

- Compile to wasm binary format. This will be needed for linking our standard library modules. Before this, we can compile to the .wat textual format and use external conversion tools.
- Hash table implementation.
- Start building out stdlib methods for core data structures as necessary.

## Week 5: 9/28 - 10/4

- Hash tables part 2.
- Implement interned symbols
- Get started on class definitions and method dispatch.

## Week 6: 10/5 - 10/11

- Continue class/method definitions and dispatch.
- Implement Javascript FFI. Javascript's Wasm API makes this pretty easy.
- Finish specifying the demo project.
- Start writing the demo project. I want to start early so I can find any blocking missing features in my compiler.
- Review my progress. Try and specify exactly what subset I want to implement by the end. Polish the architecture documentation based on the current state of the project.

## Week 7: 10/12 - 10/18

- Bug fixes! I'm sure I'll have accumulated many by now.
- Build out stdlib with features needed for the demo.

## Week 8: 10/19 - 10/25

- Bug fixes and remedial feature implementations.
- Start presentation planning and slides.

## Week 9: 10/26 - 11/1

- Bug fixes and remedial feature implementations.
- Continue presentation planning and practice.
- Start capstone writeup.

## Week 10: 10/2 - 10/8

- Bug fixes and remedial feature implementations.
- Polish the documentation.
- Present to Dr. Ladd. Incorporate feedback.
- Finish capstone writeup.

## Week 11: 10/9 - 10/15

- Bug fixes.
- Final presentation practice.
- Presentation.

## References

- [1] WC3 WebAssembly Working Group. *WebAssembly Core Specification 3.0. Draft*. May 15, 2025. URL (visited on 08/24/2025).
- [2] Andreas Rossberg et al. *GC Proposal for WebAssembly*. URL (visited on 08/24/2025).
- [3] WC3 WebAssembly Working Group. *WebAssembly Feature Status*. URL (visited on 08/24/2025).
- [4] Sritely Institute. *Hoot: Scheme on WebAssembly*. May 23, 2023. URL.
- [5] Christine Lemmer-Webber. *Directly compiling Scheme to WebAssembly: lambdas, recursion, iteration!* May 23, 2023. URL (visited on 08/24/2025).
- [6] Ruby maintainers. *ruby.wasm*. URL (visited on 08/24/2025).
- [7] Artichoke maintainers. *Build the next Ruby for Wasm with Artichoke*. URL (visited on 08/24/2025).